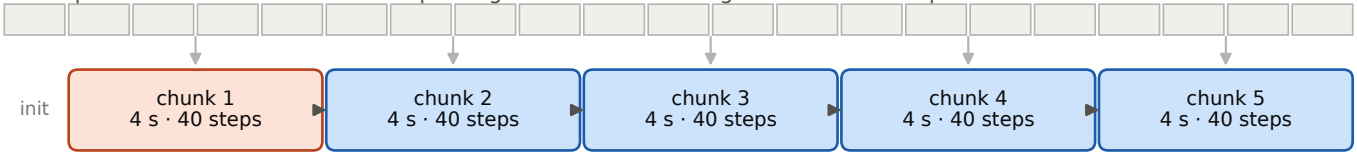

Teacher forcing: after predicting chunk k the prediction is discarded; the recorded frames and actions are registered before chunk k+1.

Backbone forwards per clip: B10 87 (video 20 + action 55 + register 12) · V0-A2 27 (5 + 10 + 12) · readout 17 (5 + 0 + 12)